

CAP 5625: Programming Assignment 2

Due on Canvas by Friday, October 27, 2023 at 11:59pm

Preliminary instructions

You may consult with other students currently taking CAP 5625 in your section at FAU on this programming assignment. If you do consult with others, then you must indicate this by providing their names with your submitted assignment. However, all analyses must be performed independently, all source code must be written independently, and all students must turn in their own independent assignment. Note that for this assignment, you may *choose* to pair up with one other student in *your section* of CAP 5625 and submit a joint assignment. If you *choose* to do this, then both your names must be associated with the assignment and you will each receive the same grade.

Though it should be unnecessary to state in a graduate class, I am reminding you that you may **not** turn in code (partial or complete) that is written or inspired by others, including code from other students, websites, past code that I release from prior assignments in this class or from past semesters in other classes I teach, or any other source that would constitute an academic integrity violation. All instances of academic integrity violations will receive a zero on the assignment and will be referred to the Department Chair and College Dean for further administrative action.

You may choose to use whatever programming language you want. However, you must provide clear instructions on how to compile and/or run your source code. I recommend using a modern language, such as Python, R, or Matlab as learning these languages can help you if you were to enter the machine learning or artificial intelligence field in the future.

All analyses performed and algorithms run must be written from scratch. That is, you may not use a library that can perform gradient descent, cross validation, ridge regression, least squares regression, optimization, etc. to successfully complete this programming assignment. The goal of this assignment is not to learn how to use particular libraries of a language, but it is to instead understand how key methods in statistical machine learning are implemented. With that stated, I will provide 5% extra credit if you additionally implement the assignment using built-in statistical or machine learning libraries (see Deliverable 6 at end of the document).

Note, credit for deliverables that request graphs, discussion of results, or specific values will not be given if the instructor must run your code to obtain these graphs, results, or specific values.

Brief overview of assignment

In this assignment you will be analyzing credit card data from $N = 400$ training observations. The goal is to fit a model that can predict credit balance based on $p = 9$ features describing an individual, which include an individual's income, credit limit, credit rating, number of credit

cards, age, education level, gender, student status, and marriage status. Specifically, you will perform a penalized (regularized) least squares fit of a linear model using ridge regression, with the model parameters obtained by batch gradient descent. The tuning parameter will be chosen using five-fold cross validation, and the best-fit model parameters will be inferred on the training dataset conditional on an optimal tuning parameter.

Data

Data for these observations are given in `Credit_N400_p9.csv`, with individuals labeled on each row (rows 2 through 401), and input features and response given on the columns (with the first row representing a header for each column). There are six quantitative features, given by columns labeled “Income”, “limit”, “Rating”, “Cards”, “Age”, and “Education”, and three qualitative features with two levels labeled “Gender”, “Student”, and “Married”.

Detailed description of the task

Recall that the task of performing a ridge regression fit to training data $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$ is to minimize the cost function

$$J(\beta, \lambda) = \sum_{i=1}^N \left(y_i - \sum_{j=1}^p x_{ij} \beta_j \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$

where y_i is a centered response and where the input p features are standardized (*i.e.*, centered and divided by their standard deviation). Moreover, recall that batch gradient descent first computes the p -dimensional gradient vector $\frac{\partial J(\beta, \lambda)}{\partial \beta}$, and then simultaneously updates each parameter k , $k = 1, 2, \dots, p$, as follows:

$$\beta_k := \beta_k - \alpha \frac{\partial}{\partial \beta_k} J(\beta, \lambda)$$

where α is the learning rate and where the partial derivative of the cost function with respect to the k th parameter is

$$\frac{\partial}{\partial \beta_k} J(\beta, \lambda) = -2 \sum_{i=1}^N x_{ik} \left(y_i - \sum_{j=1}^p x_{ij} \beta_j \right) + 2\lambda \beta_k$$

To implement this algorithm, depending on whether your chosen language can quickly compute vectorized operations, you may implement batch gradient descent using either Algorithm 1 or Algorithm 2 below (choose whichever you are more comfortable implementing). Note that in languages like R, Python, or Matlab, Algorithm 2 (which would be implemented by several nested loops) may be much slower than Algorithm 1. Note, if you are implementing Algorithm 1 using Python, use numpy arrays instead of Pandas data frames for computational speed.

You may need to explore different learning rate values to identify one that is not too large and not too small, such that it is likely for the algorithm to converge in a reasonable period of time. I

would consider a learning rate of $\alpha = 10^{-5}$, though I encourage you to explore how your model trains for smaller and larger learning rates as well. For this assignment, assume that we will reach the minimum of the cost function within a fixed number of steps, with the number of iterations being 10^5 .

Algorithm 1 (vectorized):

- Step 1.** Choose learning rate α and fix tuning parameter λ
- Step 2.** Generate N -dimensional centered response vector \mathbf{y} and $N \times p$ standardized (centered and scaled to have unit standard deviation) design matrix \mathbf{X}
- Step 3.** Randomly initialize the parameter vector $\beta = [\beta_1, \beta_2, \dots, \beta_p]$
- Step 4.** Update the parameter vector as
$$\beta := \beta - 2\alpha[\lambda\beta - \mathbf{X}^T(\mathbf{y} - \mathbf{X}\beta)]$$
- Step 5.** Repeat Step 4 for 10^5 iterations or until convergence (vector β does not change)
- Step 6.** Set the last updated parameter vector as $\hat{\beta} = [\hat{\beta}_1, \hat{\beta}_2, \dots, \hat{\beta}_p]$

Algorithm 2 (non-vectorized):

- Step 1.** Choose learning rate α and fix tuning parameter λ
- Step 2.** Generate N -dimensional centered response vector \mathbf{y} and $N \times p$ standardized (centered and scaled to have unit standard deviation) design matrix \mathbf{X}
- Step 3.** Randomly initialize the parameter vector $\beta = [\beta_1, \beta_2, \dots, \beta_p]$
- Step 4.** Create temporary parameter vector $\beta_{\text{temp}} = [\beta_1^{\text{temp}}, \beta_2^{\text{temp}}, \dots, \beta_p^{\text{temp}}]$
- Step 5.** For each $k, k = 1, 2, \dots, p$, find next value for parameter k as
$$\beta_k^{\text{temp}} := \beta_k - 2\alpha \left[\lambda\beta_k - \sum_{i=1}^N x_{ik} \left(y_i - \sum_{j=1}^p x_{ij}\beta_j \right) \right]$$
- Step 6.** Update the parameter vector as $\beta = \beta_{\text{temp}}$
- Step 7.** Repeat Steps 5 and 6 for 10^5 iterations or until convergence (vector β does not change)
- Step 8.** Set the last updated parameter vector as $\hat{\beta}$

When randomly initializing the parameter vector, I would make sure that the parameters start at small values. A good strategy here may be to randomly initialize each of the $\beta_j, j = 1, 2, \dots, p$, parameters from a uniform distribution between -1 and 1 .

Effect of tuning parameter on inferred regression coefficients

You will consider a discrete grid of seven tuning parameter values $\lambda \in \{10^{-2}, 10^{-1}, 10^0, 10^1, 10^2, 10^3, 10^4\}$ where the tuning parameter is evaluated across a wide range of values on a log scale. For each tuning parameter value, you will use gradient descent to infer the best-fit model.

Deliverable 1: Illustrate the effect of the tuning parameter on the inferred ridge regression coefficients by generating a plot (*e.g.*, using Excel, Matlab, R, etc.) of nine lines (one for each of the $p = 9$ features), with the y -axis as $\hat{\beta}_j, j = 1, 2, \dots, 9$, and the x -axis the corresponding log-scaled tuning parameter value $\log_{10}(\lambda)$ that generated the particular $\hat{\beta}_j$. Label both axes in the plot as well as provide a legend for the lines. Without the log scaling of the tuning parameter, the plot will look distorted. Note, your graph should resemble the one on Slide 16 of Module 4 lecture slides.

Choosing the best tuning parameter

You will consider a discrete grid of seven tuning parameter values $\lambda \in \{10^{-2}, 10^{-1}, 10^0, 10^1, 10^2, 10^3, 10^4\}$ where the tuning parameter is evaluated across a wide range of values on a log scale. For each tuning parameter value, perform five-fold cross validation and choose the value of λ that gives the smallest

$$CV_{(5)} = \frac{1}{5} \sum_{i=1}^5 MSE_i$$

where MSE_i is the mean squared error on the validation set of the i th-fold.

Note that during the five-fold cross validation, you will hold out one of the five sets (here 80 observations) as the Validation Set and the remaining four sets (the other 320 observations) will be used as the Training Set. On this Training Set, you will need to center the output and standardize (center and divided by the standard deviation across samples) each feature. These identical values used for centering the output and standardizing the input will need to be applied to the corresponding Validation Set, so that the Validation set is on the same scale. Because the Training Set changes based on which set is held out for validation, each of the five pairs of Training and Validation Sets will have different centering and standardization parameters.

Deliverable 2: Illustrate the effect of the tuning parameter on the cross validation error by generating a plot (*e.g.*, using Excel, Matlab, R, Python, etc.) with the y -axis as $CV_{(5)}$ error, and the x -axis the corresponding log-scaled tuning parameter value $\log_{10}(\lambda)$ that generated the particular $CV_{(5)}$ error. Label both axes in the plot. Without the log scaling of the tuning parameter, the plot will look distorted.

Deliverable 3: Indicate the value of λ that generated the smallest $CV_{(5)}$ error.

Deliverable 4: Given the optimal λ , retrain your model on the entire dataset of $N = 400$ observations and provide the estimates of the $p = 9$ best-fit model parameters.

Deliverable 5: Provide all your source code that you wrote from scratch to perform all analyses (aside from plotting scripts, which you do not need to turn in) in this assignment, along with instructions on how to compile and run your code.

Deliverable 6 (extra credit): Implement the assignment using statistical or machine learning libraries in a language of your choice. Compare the results with those obtained above, and provide a discussion as to why you believe your results are different if you found them to be different. This is worth up to 5% additional credit, which would allow you to get up to 105% out of 100 for this assignment.